

AI-POWERED KNOWLEDGE INTERACTION SYSTEM

Assistant Professor D. Mulli Rogers

UG Students K.Ebenazer, S.Arish Doni, V.Bebin, RM.Rahul

Department Of Artificial Intelligence & DataScience
St.Xavier's Catholic College of Engineering,
Chunkankadai-629003, Tamil Nadu

Abstract:- With the rapid increase in digital documents, extracting useful information from PDF files has become challenging in recent years. Due to this, intelligent document analysis methods using Artificial Intelligence are widely adopted. One of the important applications of Natural Language Processing is the Retrieval-Augmented Generation (RAG) based chatbot. It reduces manual effort, improves information accessibility, and enhances decision making. RAG-based chatbot is an automated and intelligent system implemented using large language models and vector databases. In this project, PDF document processing and semantic search are studied and implemented based on real-world document analysis requirements. The system extracts text from uploaded PDF files, stores the content using embeddings, and retrieves relevant information based on user queries. The major objective is to provide accurate, context-aware answers from documents and to identify financial data such as revenue, expenses, and profits. By detecting financial information, the system generates visual representations such as bar charts and pie charts, enabling better data understanding and analysis.

Keywords— RAG chatbot, Financial analysis, PDF retrieval, Vector database, Data visualization.

I.INTRODUCTION

The word 'Information' springs from the Latin word 'Informare' means to give form or knowledge and 'Technology' means application. It means the science and art of processing data and documents for meaningful purpose. Information technology is an art of extracting knowledge from digital documents for the use of mankind. Information management is the mile stone within the history of human civilization, thanks to information systems man managed data efficiently. Information processing is one amongst the oldest and prime activities of the human being. It's remained a vital source of knowledge.

In spite of growing digitalization and automation within the world, nearly one half of useful information still remains stored in unstructured documents. In developing Countries document based systems has been a major source of information and it's contributed to the economy. The fundamental aim of document processing is to extract accurate and relevant information and to help users for decision making by improving accessibility and understanding. Information systems are a backbone of modern economy. In India about sixty four percent of organizations are dependent on documents for their daily operations. The document processing activities in the world are closely controlled by computational factors.

Document analysis is not an exception for this, today systems are facing two main problem concerned with documents. The first is meeting the increasing demand of information access and other is processing large volume of documents for ever increasing data and the second is uneven extraction of knowledge and changing pattern of document usage. Organizations tried to become efficient in document handling through advanced technologies. After digital revolution by taking systematic efforts due to the unique

The study of documents and information from the computational point of view gained more importance after 2000. At the beginning of 2010 and later on artificial intelligence brought of remarkable change in the field of information processing, due to this systems become not only efficient in retrieval but it could also generate intelligent responses. The process of document analysis is not properly channelized because of unstructured data, unavailability of context awareness and unbalance processing of information. Intelligent systems succeed only with proper data retrieval. In spite of lot of efforts by developers, users could not get accurate answers. This creates a large gap between data and knowledge and imbalance is created. To reduce this gap. Systematic design is required for this purpose it is necessary to have the detailed information extraction.

Many organizations like India, majority of the operations depend on documents, and its knowledge base comes from documents. In spite of this and even the modern technology is found everywhere, the document processing area is following the old conventional methods. Users still resort to traditional search methods, manual reading of documents, unscientific systems of analysis. The data is unstructured, and unevenness of information throughout the document poses a major

problem. All this leads to inaccurate results and low productivity. The implementation of intelligent methods in the field of document analysis can bring about radical changes in the accuracy of information, due to improved efficiency in retrieval techniques. Of the various advantages that RAG brings to the table, its ability to integrate retrieval and generation is absolutely ground- breaking. Mostly, we come across ideas that suggest a retrieval system that collects data from the various documents present in the database

and sends the data to the language model. This method focuses on studying the document content to improve response accuracy. But it turns out, retrieving information alone are never adequate to generate meaningful answers since a lot of other factors have a role to play.

II. LITERATURE SURVEY

1.Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks

Lewis et al. (2020) proposed Retrieval-Augmented Generation (RAG) for improving knowledge-intensive natural language processing tasks. The system integrates a dense document retriever with a sequence-to-sequence generative model to enhance contextual understanding. By retrieving relevant documents before generating responses, the model improves factual accuracy and reduces hallucination. The proposed system was evaluated on open-domain question answering datasets and achieved better performance than traditional parametric language models. However, the implementation focused mainly on general knowledge datasets and did not specifically address financial document processing, structured PDF extraction, or visualization of financial metrics for analytical decision-making.

2. Language Models are Few-Shot Learners

Brown et al. (2020) introduced GPT-3, a large-scale transformer-based language model capable of performing multiple natural language processing tasks using few-shot learning. The model demonstrated strong performance in text generation, summarization, translation, and question answering. It relies heavily on pretrained knowledge and large-scale training data. Although highly effective for conversational applications, the system lacks a built-in mechanism for retrieving real-time information from external financial reports. This limitation may lead to outdated or imprecise responses when applied to financial document analysis tasks.

3. Dense Passage Retrieval for Open-Domain Question Answering

Karpukhin et al. (2020) proposed Dense Passage Retrieval (DPR), which uses dual encoders to map queries and documents into the same embedding space. Unlike traditional TF-IDF-based retrieval methods, DPR captures semantic meaning through contextual embeddings. The system demonstrated superior retrieval accuracy on open-domain QA datasets such as Natural Questions and TriviaQA. DPR significantly improves recall rates and reduces irrelevant document retrieval. However, it focuses solely on retrieval and does not integrate generative response mechanisms or analytical processing. Additionally, the study does not explore financial domain

adaptation or structured document analysis. Thus, integrating DPR into a financial RAG-based chatbot remains an area of improvement.

4. FinBERT: Financial Sentiment Analysis Using BERT

Araci (2019) developed FinBERT, a domain-specific BERT model fine-tuned on financial text corpora. The objective was to enhance sentiment analysis accuracy in financial news and reports. FinBERT achieved improved classification performance compared to general-purpose BERT models. The model is particularly useful for identifying market sentiment and investor behavior. However, it focuses only on sentiment polarity detection and does not perform document retrieval, financial metric extraction, or conversational interaction. Moreover, it does not generate visual summaries or perform ratio analysis. Hence, while FinBERT improves financial text understanding, it does not provide an integrated analytical chatbot framework.

5. FAISS: Efficient Similarity Search and Clustering of Dense Vectors

Johnson et al. (2017) introduced FAISS, a library developed by Facebook AI Research for efficient similarity search in large-scale vector databases. FAISS supports GPU acceleration and approximate nearest neighbor search, making it suitable for handling millions of embeddings efficiently.

This system plays a crucial role in retrieval-based AI applications, especially in semantic search systems. However, FAISS acts only as an indexing and retrieval engine. It does not provide contextual reasoning, financial calculations, or visualization capabilities. Therefore, additional components must be integrated to build a complete financial intelligence chatbot system.

6. Attention Is All You Need

Vaswani et al. (2017) proposed the Transformer architecture, introducing self-attention mechanisms that replaced traditional recurrent neural networks. The model improved parallel processing efficiency and contextual representation learning. Transformers form the backbone of modern NLP systems including BERT, GPT, and RAG. Despite its revolutionary impact, the original transformer model does not include external document retrieval or domain-specific financial computation modules. It processes input sequences but lacks mechanisms for structured PDF parsing and graphical analytics integration.

7. A Survey of Chatbot Architectures and Applications

Adamopoulou and Moussiades (2020) presented a comprehensive survey of chatbot systems, analyzing their evolution, design methodologies, and real-world applications across various industries. The study categorizes

chatbots into three major types: rule-based chatbots, retrieval-based chatbots, and generative chatbots. Rule-based systems operate on predefined decision trees and pattern-matching techniques such as AIML (Artificial Intelligence Markup Language). While these systems are simple and easy to deploy, they lack flexibility and fail when user queries deviate from predefined rules.

Retrieval-based chatbots use machine learning algorithms to select the most appropriate response from a predefined dataset. These systems improve adaptability compared to rule-based bots but are still limited by the scope of stored responses. Generative chatbots, on the other hand, utilize deep learning models such as sequence-to-sequence architectures and transformers to dynamically generate responses. The study emphasizes that transformer-based models significantly enhance conversational fluency, contextual understanding, and scalability.

8. Data Visualization for Business Intelligence

Few (2006) highlighted the critical importance of data visualization in business intelligence systems. The study explains how visual representations such as bar charts, line graphs, pie charts, scatter plots, and dashboards transform complex datasets into understandable insights. Visualization plays a vital role in identifying trends, anomalies, patterns, and performance indicators within large financial datasets.

The research emphasizes principles of effective visual communication, including clarity, minimalism, accuracy, and appropriate selection of chart types. For example, time-series financial data is best represented using line charts, while comparative performance metrics are better suited for bar graphs. Proper visualization reduces cognitive load and improves executive decision-making efficiency.

In financial domains, visualization supports revenue analysis, profit margin comparison, cash flow monitoring, and investment trend tracking. Modern business intelligence tools integrate interactive dashboards that allow filtering, zooming, and real-time updates. However, most traditional visualization systems depend heavily on structured datasets stored in databases or spreadsheets. They require manual preprocessing of financial data before generating graphical outputs.

9. LayoutLM: Pre-training of Text and Layout for Document Image Understanding

Xu et al. (2020) proposed LayoutLM, a pre-training model designed for document image understanding tasks. Unlike traditional NLP models that process only textual information, LayoutLM incorporates both textual content and layout (spatial position) information from documents. The objective of this research was to improve performance in document classification, information extraction, and form understanding tasks by combining text embeddings with 2D positional embeddings.

The model architecture is based on the Transformer encoder and extends BERT by integrating spatial layout features. Each token embedding includes three components: word embedding, positional embedding, and layout embedding. By capturing the relative positioning of text within documents, LayoutLM improves contextual understanding in structured documents such as invoices, receipts, forms, and financial statements.

Experimental results demonstrated significant performance improvements on benchmark datasets like FUNSD (Form Understanding in Noisy Scanned Documents). The study proved that combining visual and textual information enhances document intelligence systems. LayoutLM and its later versions (LayoutLMv2 and LayoutLMv3) further improved document image processing by integrating visual backbone networks and multi-modal pretraining.


III. EXISTING SYSTEM

Existing financial analysis systems mainly rely on manual inspection of financial reports in PDF format. Analysts must read lengthy annual reports, balance sheets, and profit and loss statements to extract meaningful insights. Traditional tools such as spreadsheets and static dashboards provide visualization support but do not intelligently retrieve context-specific information from large unstructured documents. Financial data is often distributed across multiple PDFs, making it difficult to search and analyze efficiently. Conventional chatbot systems are limited to predefined datasets and cannot dynamically retrieve information from uploaded documents. These systems lack advanced retrieval mechanisms, resulting in generic or incomplete responses. Existing solutions often require technical expertise to generate financial ratios and performance metrics. There is no integrated framework that combines document retrieval, intelligent querying, financial computation, and visualization in a unified manner, leading to time-consuming analysis and reduced productivity.

IV. PROPOSED SYSTEM

This project presents a proposed model for an Intelligent Financial Document Analysis and Monitoring System using Retrieval-Augmented Generation (RAG) architecture. The system is designed to process financial PDF documents, extract relevant data such as revenue, expenses, assets, liabilities, and profit margins, and provide real-time analytical insights through an interactive chatbot interface. It enables users to upload financial reports and receive contextual responses along with graphical representations through web and mobile applications. The AI-based monitoring system helps users understand financial performance trends and supports decision-making remotely from anywhere, anytime. The integration of semantic search with generative AI improves the accuracy of responses and reduces misinformation. The proposed system utilizes document embedding techniques and vector databases to

store and retrieve relevant financial information efficiently. The architecture consists of a PDF parser, embedding generator, vector indexing module, retrieval engine, and transformer-based



language model for contextual response generation. The system extracts structured financial metrics from unstructured reports and performs automated ratio analysis including profitability ratio, liquidity ratio, and growth comparison. The chatbot provides responses in natural language format and can also generate charts such as bar graphs and line graphs for better visualization of financial trends.

The backend is developed using a scalable server framework integrated with a vector database and API services. The frontend interface is designed for both web and mobile platforms, ensuring accessibility and ease of use. The system supports secure document handling and real-time interaction. By combining document intelligence, retrieval mechanisms, financial computation, and data visualization, the proposed model significantly reduces manual analysis effort, minimizes human error, and enhances analytical efficiency. The system is flexible, scalable, and can be deployed in various financial environments such as corporate analysis, auditing, and investment monitoring.

1.Features:

- No complex configuration required, simple and user-friendly interface
- Upload financial PDF documents into the system, ask queries through chatbot and view analytical results instantly
- Supports secure document processing with scalable cloud-based architecture
- Lightweight web and mobile application with fast response time
- Designed for financial analysts, students, auditors, and business professionals
- Eliminates manual data extraction from lengthy financial reports
- Prevents calculation errors through automated financial ratio computation
- Easy to operate, just upload document, ask question, and receive response with visualization

2.Applications:

- Financial report analysis
- Investment performance monitoring
- Audit and compliance support
- Business and academic research

V. CONCLUSION

The financial analysis domain plays a crucial role in supporting strategic decision making across organizations. With the rapid growth of digital financial documentation, traditional analysis methods are becoming time-consuming and inefficient. The increasing volume of PDF-based reports and unstructured financial data has created challenges in extracting accurate insights and performing timely evaluations. Conventional systems lack intelligent retrieval capabilities and require significant manual effort, which may lead to reduced productivity and delayed decision processes.

The proposed Retrieval Augmented Generation based financial chatbot provides an effective solution to these challenges by integrating document retrieval, automated analysis, and data visualization within a unified framework. By leveraging modern embedding models and vector databases, the system enhances the accuracy and relevance of financial insights. The framework improves analytical efficiency, reduces manual workload, and enables scalable deployment across various financial domains. With proper implementation and future enhancements, the system can significantly contribute to intelligent financial management and data-driven decision making.

Acknowledgment

The satisfaction that accompanies the successful completion of this project would be incomplete without the support of our management. I express my sincere gratitude and respect to all those who guided and inspired me throughout this research work.

REFERENCES

[1] Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich K ttler, Mike Lewis, Wen-tau Yih, Tim Rockt schel, Sebastian Riedel, and Douwe Kiela (2020)

[2] Language Models are Few-Shot Learners

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D. Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam

McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei (2020)

[3] Dense Passage Retrieval for Open-Domain Question Answering

Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wentaoyih (2020)

[4] FinBERT: Financial Sentiment Analysis Using BERT

Dogu Araci (2019)

[5] FAISS: Efficient Similarity Search and Clustering of Dense Vectors

Jeff Johnson, Matthijs Douze, and Hervé Jégou (2017)

[6] Attention Is All You Need

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin (2017)

[7] A Survey of Chatbot Architectures and Applications

Eirini Adamopoulou and Lefteris Moussiades(2020)

[8] Data Visualization for Business Intelligence

Stephen Few (2006)

[9] LayoutLM: Pre-training of Text and Layout for Document Image Understanding

Yiheng Xu, Minghao Li, Lei Cui, Shaohan Huang, Furu Wei, and Ming Zhou (2020)

